



Data Structures with C++ : CS189

Lecture 8-1: Balancing Trees

Recap

- A binary search tree speeds up lookups by keeping all data less than a node to the left and all data greater to the right
- This is $\log n$ because every decision you make cuts half of the tree off from being considered
- BST is what's inside a Set, but a BST is not itself an STL container

Balancing

- If we took the tree we have so far and entered a set of data that was in order, you wouldn't have a tree
 - 1 2 3 4 5 6 7 8 9 makes a straight line down the right side
- That's $O(n)$ list-searching. We want $O(\log n)$ so we need to do something to change the tree
- To "balance" a tree is to try to make it as even as possible

Types of Balancing

- There are two different times you could balance a tree
 - As you add or remove data
 - All at once when the tree's not being used
- There are two paradigms for how to balance
 - Fast but okay
 - More overhead but better (more tightly packed)
- All of the choices are right, they just depend on the program and the data

AVL (p 257)

- This is a faster-with-overhead method that has the benefit of being easier to write
 - I'll explain when I talk about the homework why this isn't scary
- The idea is that if the height* of the two sides of a node differ by two, we will even the two sides out
 - *Height = Greatest number of nodes to get to a leaf
- Since we do this on every add and delete, the tree is always kept *mostly* balanced

Out of Balance

- This is another example of this class being impossible without drawing pictures
- Picture A is in balance.
 - The number in the circle is the difference in height between the sides
 - "h" means "0 or more nodes that I don't care about" They aren't part of the rotate, except to note that their biggest height is h
- In picture B, a node has been added to the rightmost subtree, and now P is out of balance
 - Left is h, right is $h + 2$ ($h + 1 + Q$)

Rotation (p 257, fig 6.41)

- Picture c is the result of balancing
- We take P, Q, and P's left and rotate the three nodes counter-clockwise
 - Note the tricky move Q's left makes
- We have maintained BST-ness as left data is still less and right data is still greater
 - That middle tree went from "Greater than P and less than Q" to "Less than Q and greater than P"
- You can see this code in the files provided
 - The assignment is going to be optimizing that file, not writing it from scratch

Double Rotation (p 258, fig 6.42)

- That rotation works if the node added was right-right or left-left (outside)
- If the add went right-left or left-right (inside) then we need to rotate twice
- (a) is in balance, and (b)'s right-left add put it out of balance
- (c) is doing nothing. It is just observing that Q's left could be called R
- (d) Q and R rotate clockwise
- (e) P and R rotate counter-clockwise
- And again, I wrote this. I just wrote it badly for you.

My Implementation

- In my code, you will see that I recursively balance every node in the tree at every add.
- You will also see that I recursively set the heights of every node twice for every add
- But this algorithm never forks
 - If you can use a loop, use a loop
 - The path of adding a node is a straight line
- A rotate on an add means it is done, but a delete always has to check all the way up

Your Task

I also summarize this in the assignment.

1. Keep track of the path you take to do your Add or Remove
2. Only update the heights of those nodes
 - a. Remember that two-child delete needs to consider the path to the actual deleted node too
3. Only check for balance on those nodes
4. If you make a rotation, update the heights of the nodes involved
5. If you make a rotation on an add, you are done
 - a. Remove could make a chain of rotates



End

Draw this out, and ask questions I can draw on the board